

Visual Language Model for
Few-Shot Learning
by DeepMind



Flamingo

Мотивация

Картинка

Текст

Vision Language

Задачи

1. Visual question-answering
2. Captioning
3. Visual dialogue


Еще мотивация

- Хотим генеративную модель,
понимающую и текст и картинки
- Хотим как в GPT-3 уметь in-context
few-shot

Как бы мы это делали?


Вход:

<image> Что изображено на этой картинке?



Embeddings:


img_1, word_1, word_2, ... word_k



Трансформер:


Большой желтый кот

Архитектура



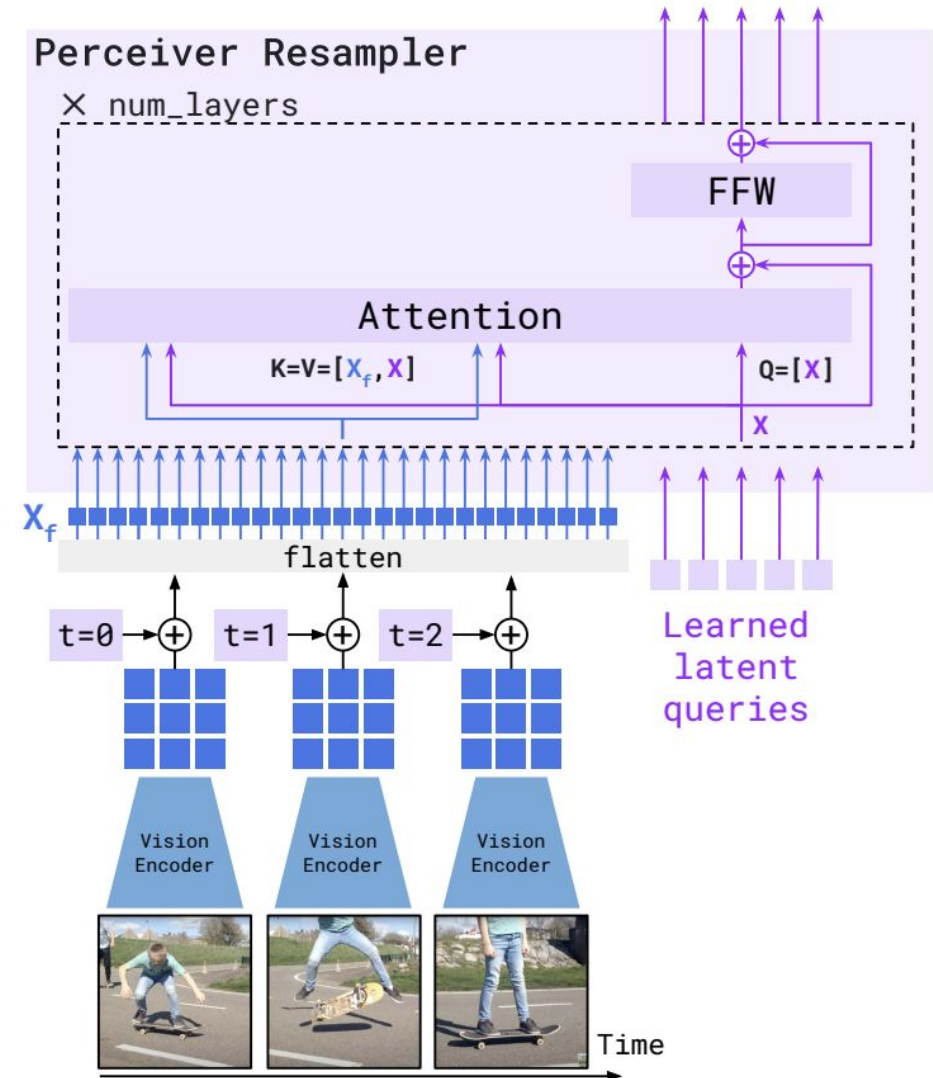
Vision Encoder

- Взяли свой продвинутый ResNet (NFNet)
- Обучили его как CLIP
- Заморозили



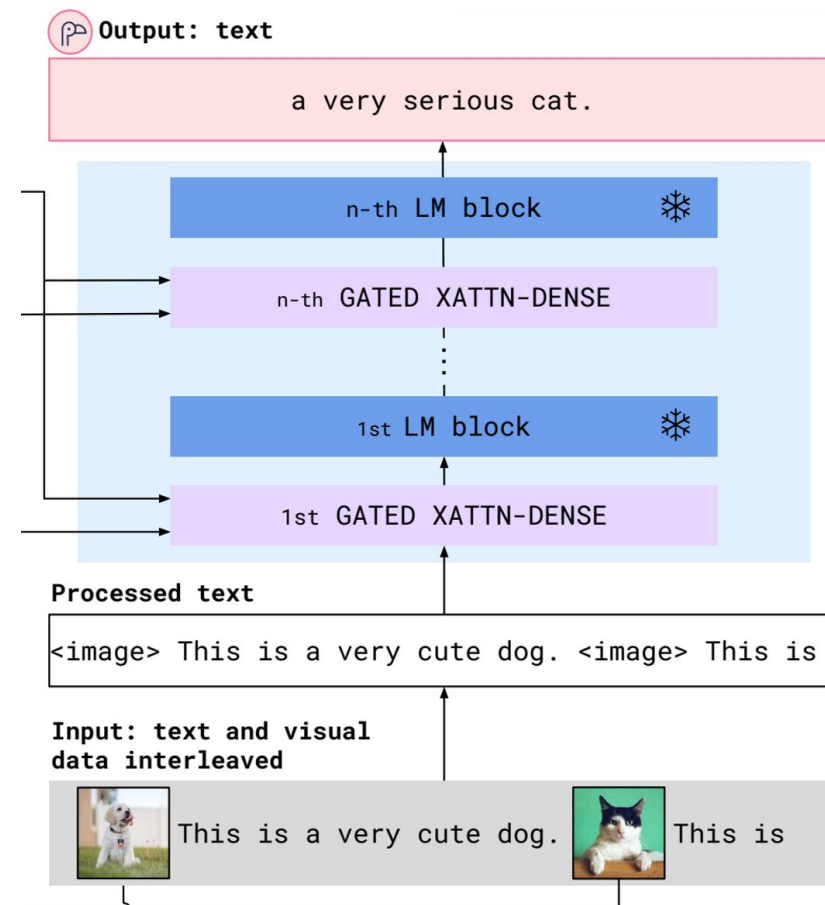
Perceiver Resampler

- Принимает на вход **feature-map** произвольной длины
- Выдает **feature-map** фиксированной длины

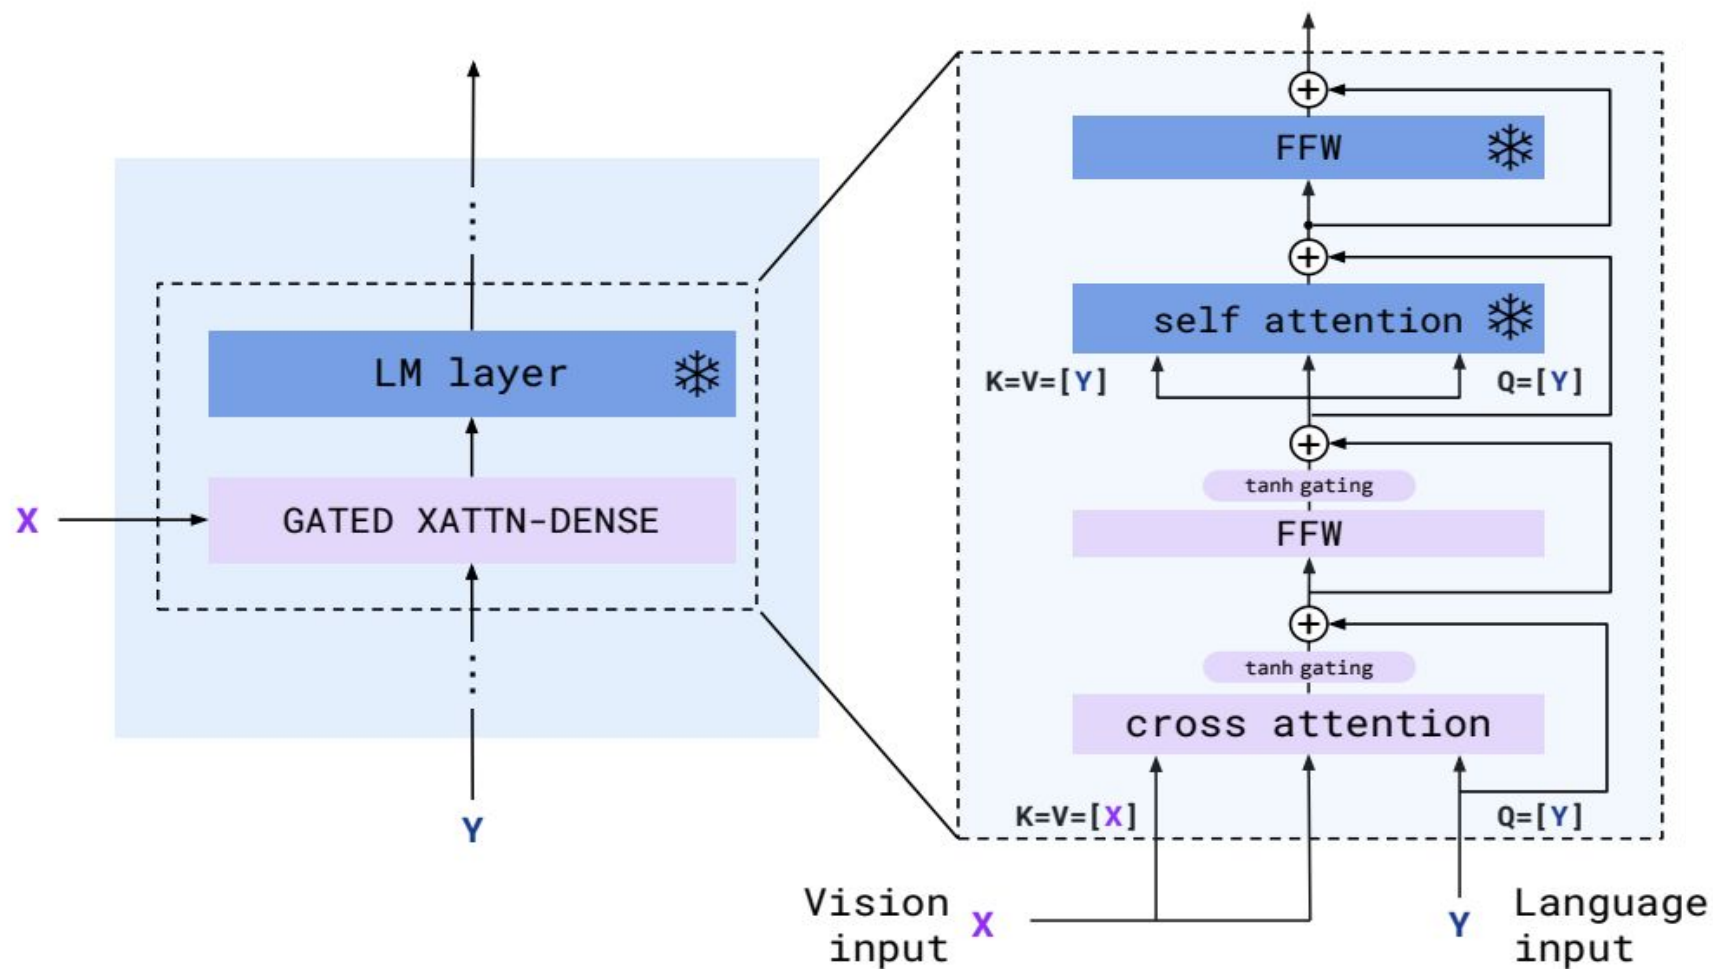


Language Encoder

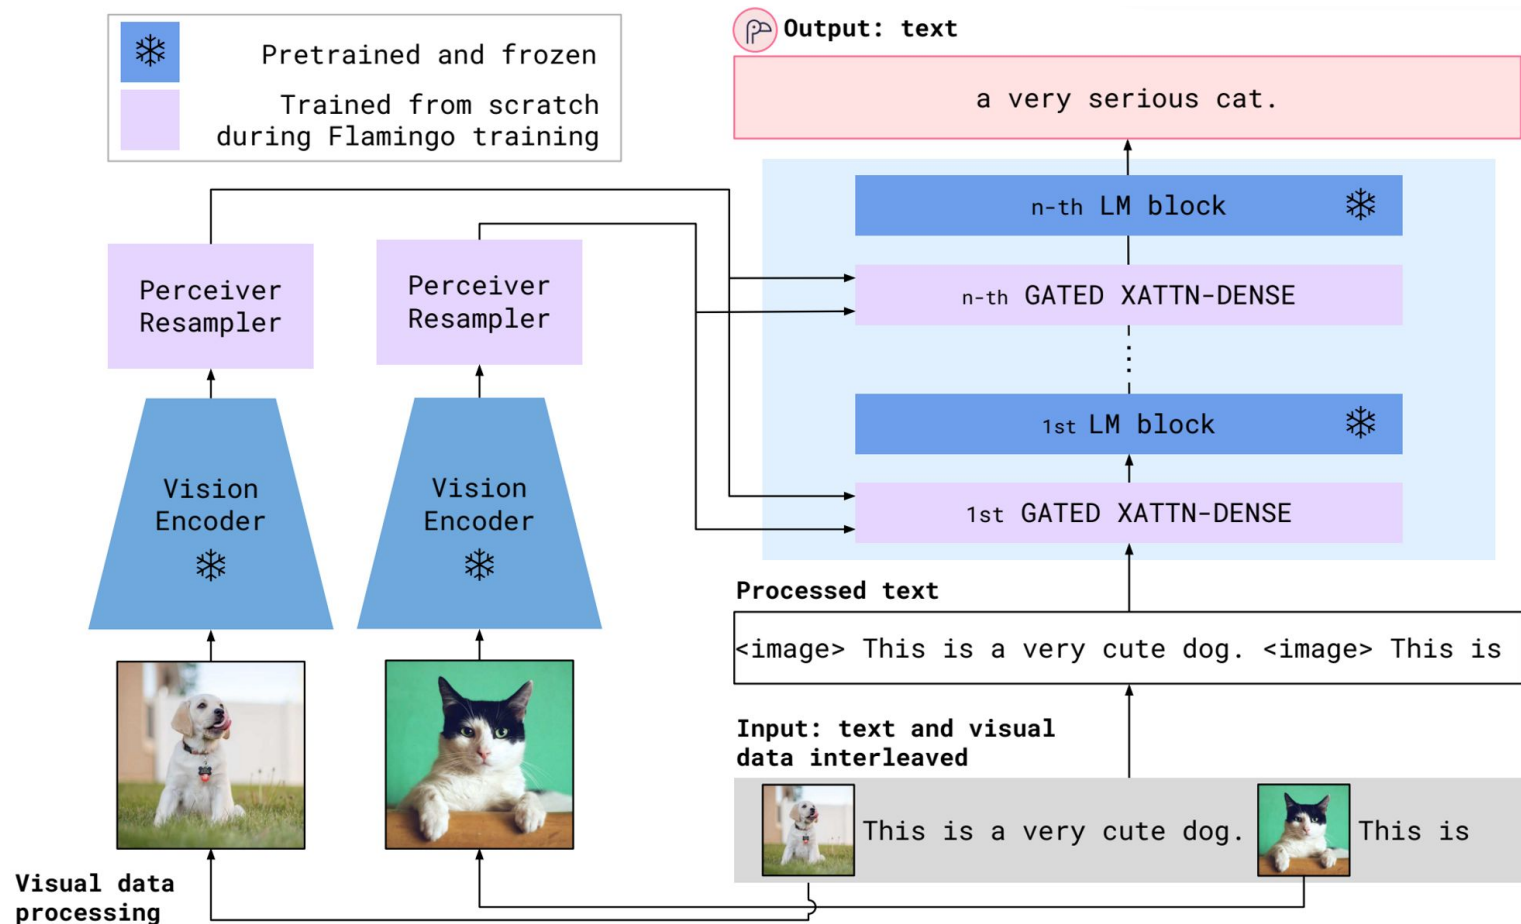
- GPT-like моделька Chinchilla (70B)
- Никак не дообучали
- Заморозили
- Добавили обучаемые слои Cross-Attention



Cross Attention



Опять архитектура



Данные

1. Веб-страницы с картинками
2. Картинки с описаниями
3. Короткие видеосюжеты с описаниями

Взвешиваем лосс

Эксперименты!




Figure 2: Flamingo results overview. *Left:* Our largest model, dubbed *Flamingo*, outperforms state-of-the-art fine-tuned models on 6 of the 16 tasks we consider with no fine-tuning. For the 9 tasks with published few-shot results, *Flamingo* sets the new few-shot state of the art. *Note:* We omit RareAct, our 16th benchmark, as it is a zero-shot benchmark with no available fine-tuned results to compare to. *Right:* Flamingo performance improves with model size and number of shots.

Больше экспериментов!



Method	VQAV2		COCO	VATEX	VizWiz		MSRVTTQA	VisDial		YouCook2	TextVQA		HatefulMemes
	test-dev	test-std			test-dev	test-std		valid	test-std		valid	test-std	
 32 shots	67.6	-	113.8	65.1	49.8	-	31.0	56.8	-	86.8	36.0	-	70.0
 Fine-tuned	82.0	82.1	138.1	84.2	65.7	65.4	47.4	61.8	59.7	118.6	57.1	54.1	86.6
SotA	81.3 [†]	81.3 [†]	149.6[†]	81.4 [†]	57.2 [†]	60.6 [†]	46.8	75.2	75.4[†]	138.7	54.7	73.7	84.6 [†]
	[133]	[133]	[119]	[153]	[65]	[65]	[51]	[79]	[123]	[132]	[137]	[84]	[152]

Table 2: Comparison to SotA when fine-tuning *Flamingo*. We fine-tune *Flamingo* on all nine tasks where *Flamingo* does not achieve SotA with few-shot learning. *Flamingo* sets a new SotA on five of them, outperforming methods (marked with †) that use tricks such as model ensembling or domain-specific metric optimisation (e.g., CIDEr optimisation).

Конец


Conclusion. We proposed Flamingo, a general-purpose family of models that can be applied to image and video tasks with minimal task-specific training data. We also qualitatively explored interactive abilities of *Flamingo* such as “chatting” with the model, demonstrating flexibility beyond traditional vision benchmarks. Our results suggest that connecting pre-trained large language models with powerful visual models is an important step towards general-purpose visual understanding.

Источники

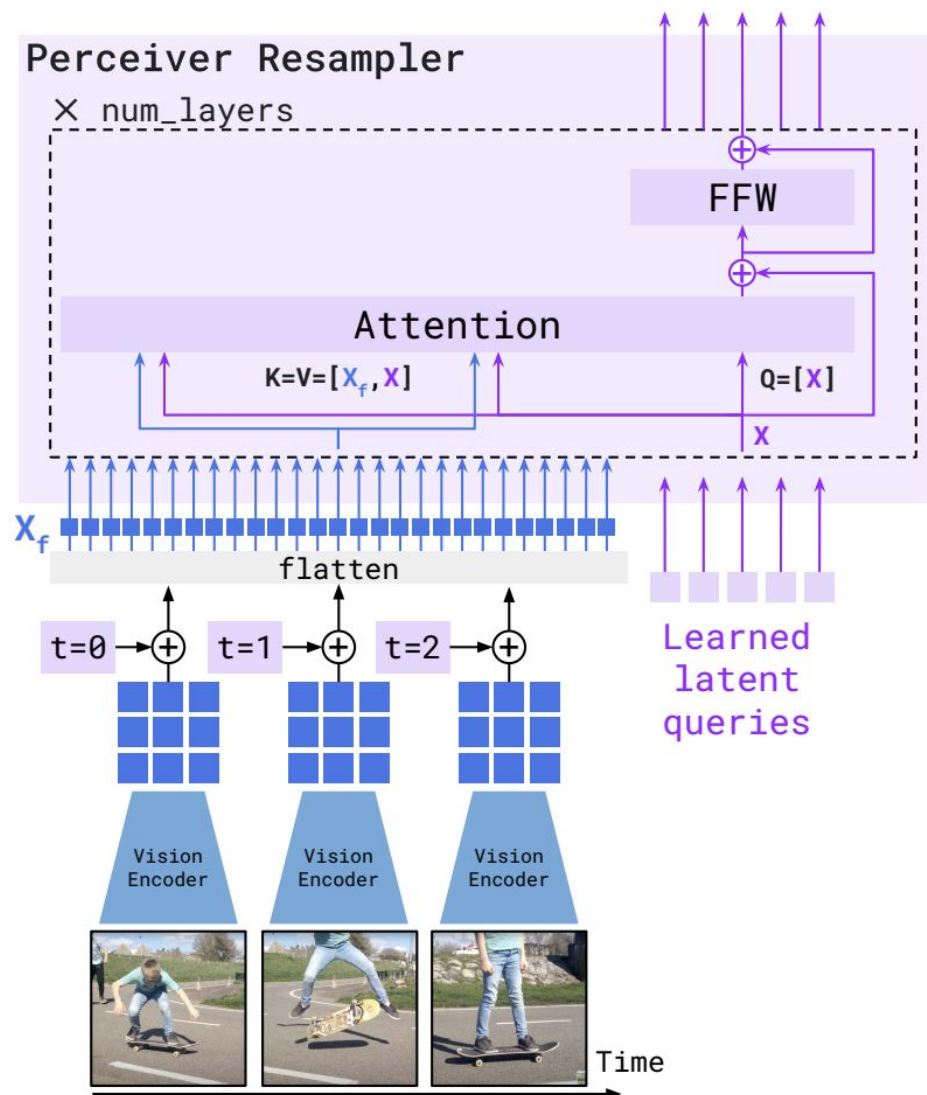
1. DeepMind [Paper](#)
2. DeepMind [Blogpost](#)
3. Small [Overview](#)
- 4.

Для справки

NFNet



Resampler



```
def perceiver_resampler(
    x_f, # The [T, S, d] visual features (T=time, S=space)
    time_embeddings, # The [T, 1, d] time pos embeddings.
    x, # R learned latents of shape [R, d]
    num_layers, # Number of layers
):
    """The Perceiver Resampler model."""

    # Add the time position embeddings and flatten.
    x_f = x_f + time_embeddings
    x_f = flatten(x_f) # [T, S, d] -> [T * S, d]
    # Apply the Perceiver Resampler layers.
    for i in range(num_layers):
        # Attention.
        x = x + attention_i(q=x, kv=concat([x_f, x]))
        # Feed forward.
        x = x + ffw_i(x)
    return x
```

Masking details

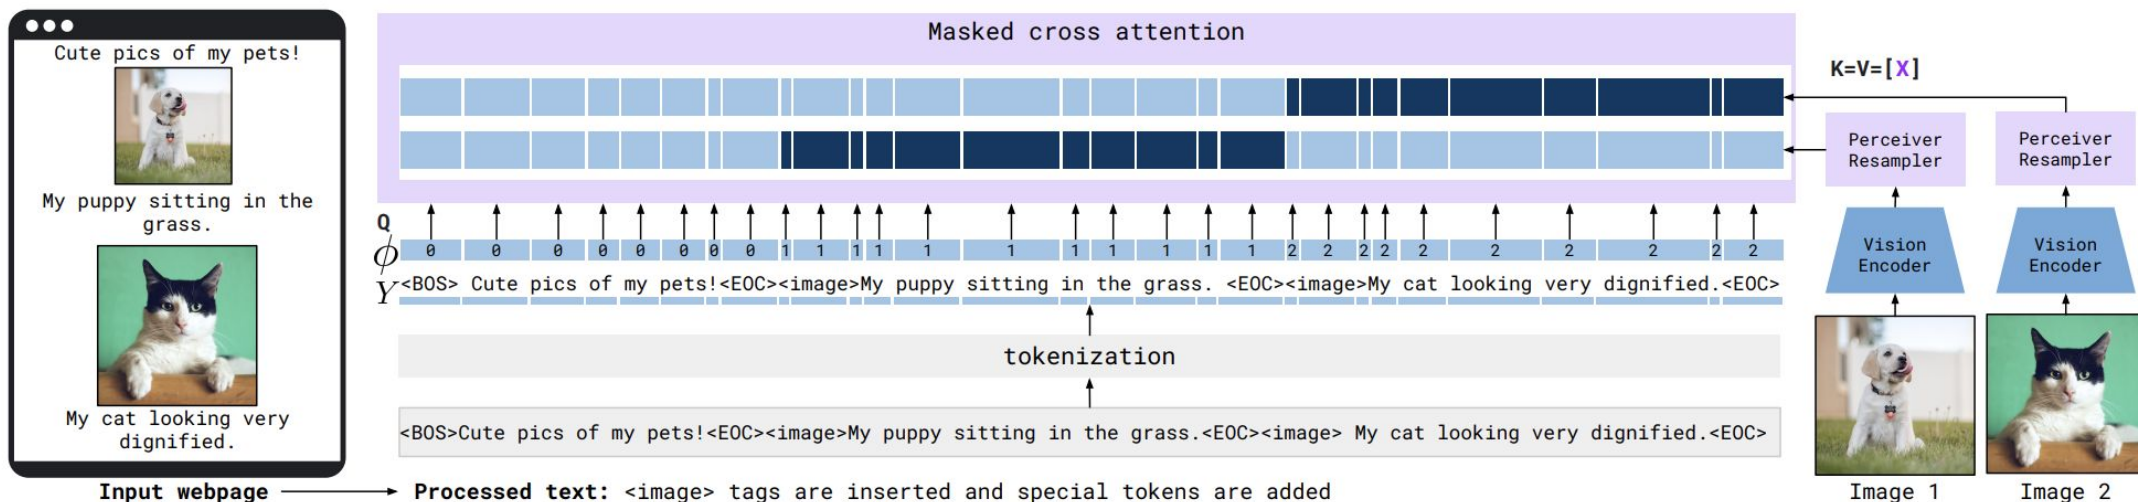


Figure 7: Interleaved visual data and text support. Given text interleaved with images/videos, e.g. coming from a webpage, we first process the text by inserting `<image>` tags at the locations of the visual data in the text as well as special tokens (`<BOS>` for “beginning of sequence” or `<EOC>` for “end of chunk”). Images are processed independently by the Vision Encoder and Perceiver Resampler to extract visual tokens. At a given text token, the model only cross-attends to the visual tokens corresponding to the last preceding image/video. ϕ indicates which image/video a text token can attend or 0 when no image/video is preceding. In practice, this selective cross-attention is achieved through masking – illustrated here with the dark blue entries (unmasked/visible) and light blue entries (masked).