Jukebox: A Generative Model for Music

Speaker: Petr Grinberg

Recall: VQ-VAE

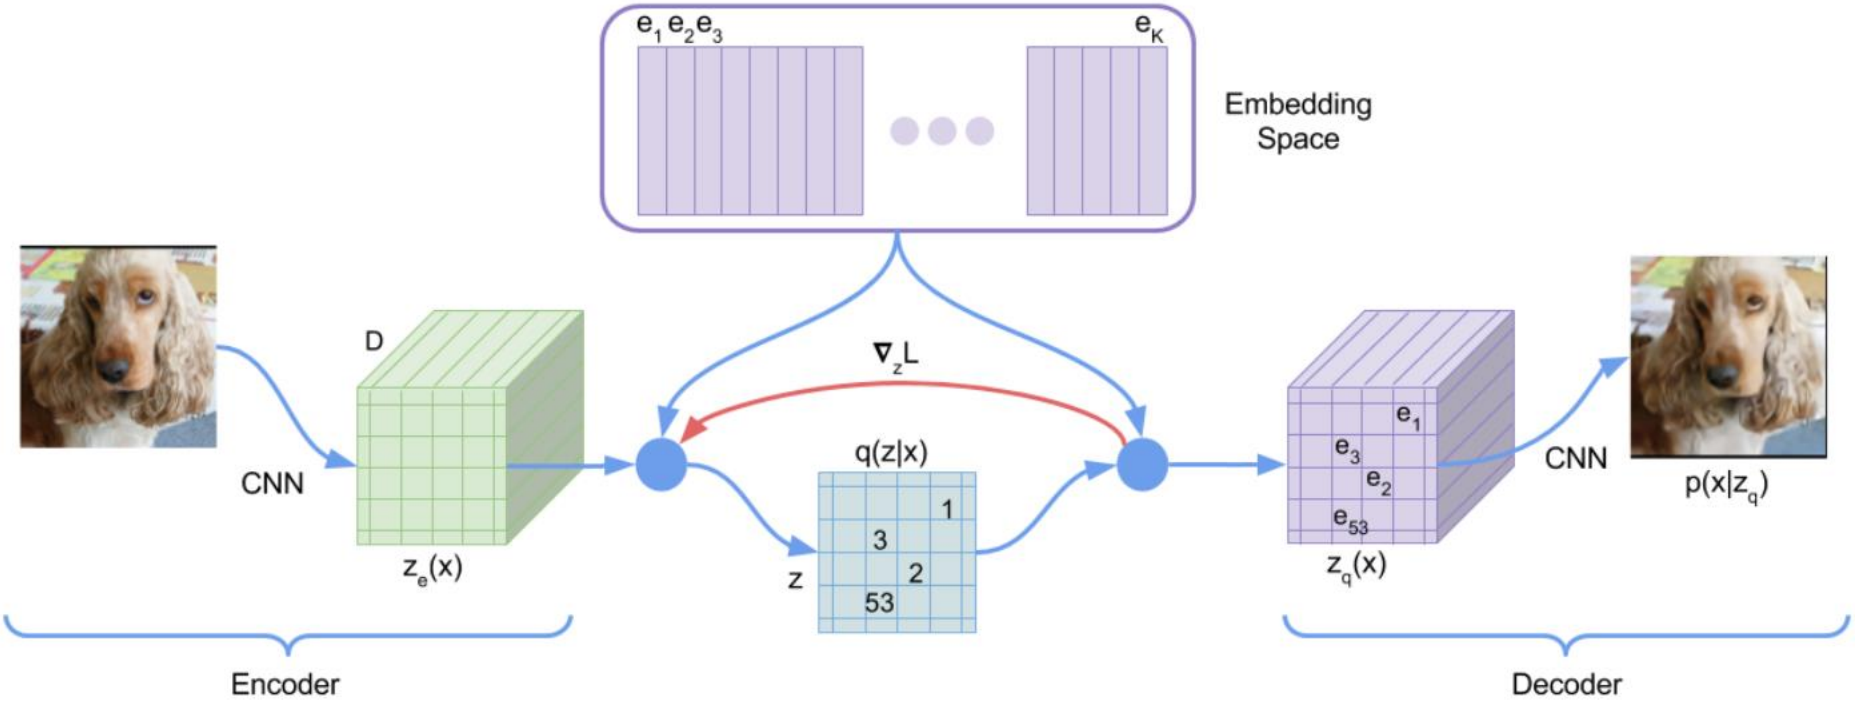

New pictures can be generated from autoregressive distribution over z, p(z)

Recall: VQ-VAE, Training Loss

$$\mathcal{L} = \mathcal{L}_{\text{recons}} + \mathcal{L}_{\text{codebook}} + \beta\mathcal{L}_{\text{commit}}$$

$$\mathcal{L}_{\text{recons}} = \frac{1}{T}\sum_{t} \|\mathbf{x}_t - D(\mathbf{e}_{z_t})\|_2^2$$

$$\mathcal{L}_{\text{codebook}} = \frac{1}{S}\sum_{s} \|\text{sg}\left[\mathbf{h}_s\right] - \mathbf{e}_{z_s}\|_2^2$$

$$\mathcal{L}_{\text{commit}} = \frac{1}{S}\sum_{s} \|\mathbf{h}_s - \text{sg}\left[\mathbf{e}_{z_s}\right]\|_2^2$$

Reconstruction Loss can vary, this one will be used in Jukebox

Recall: VQ-VAE2 (Hierarchical)

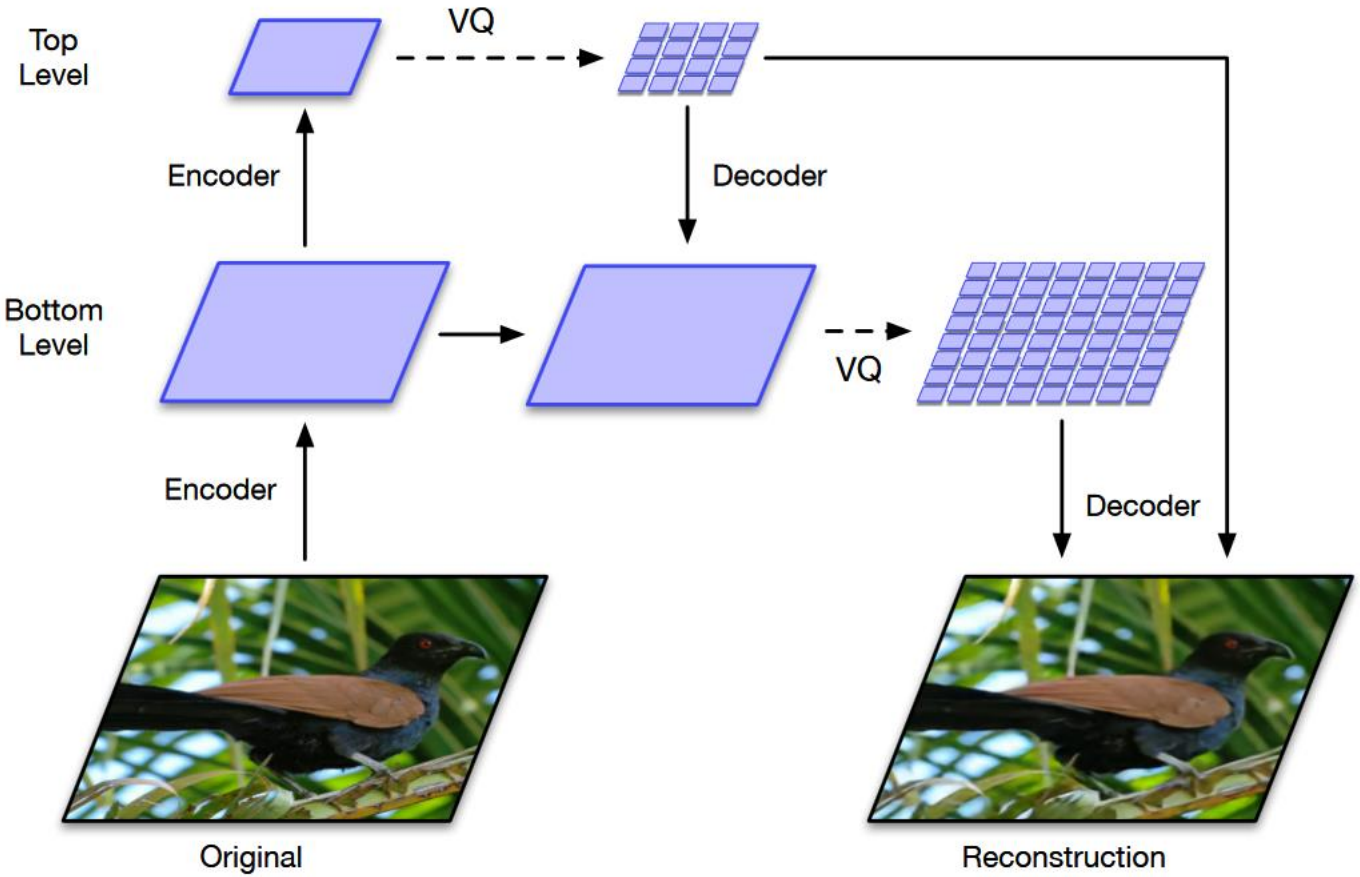

Bottom Level: local features, Top Level: high-level semantics

Jukebox: VQ-VAE

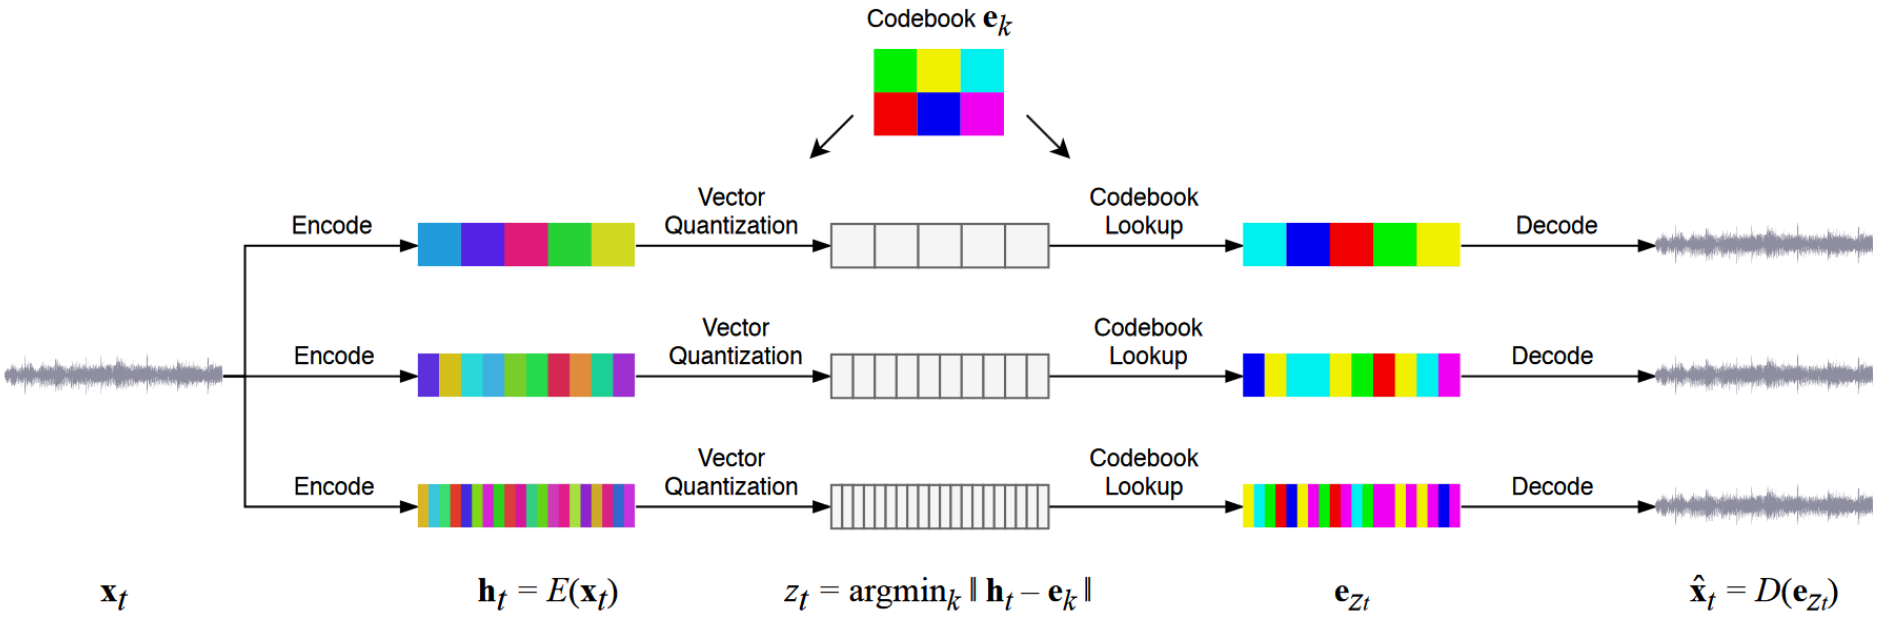

Jukebox: VQ-VAE

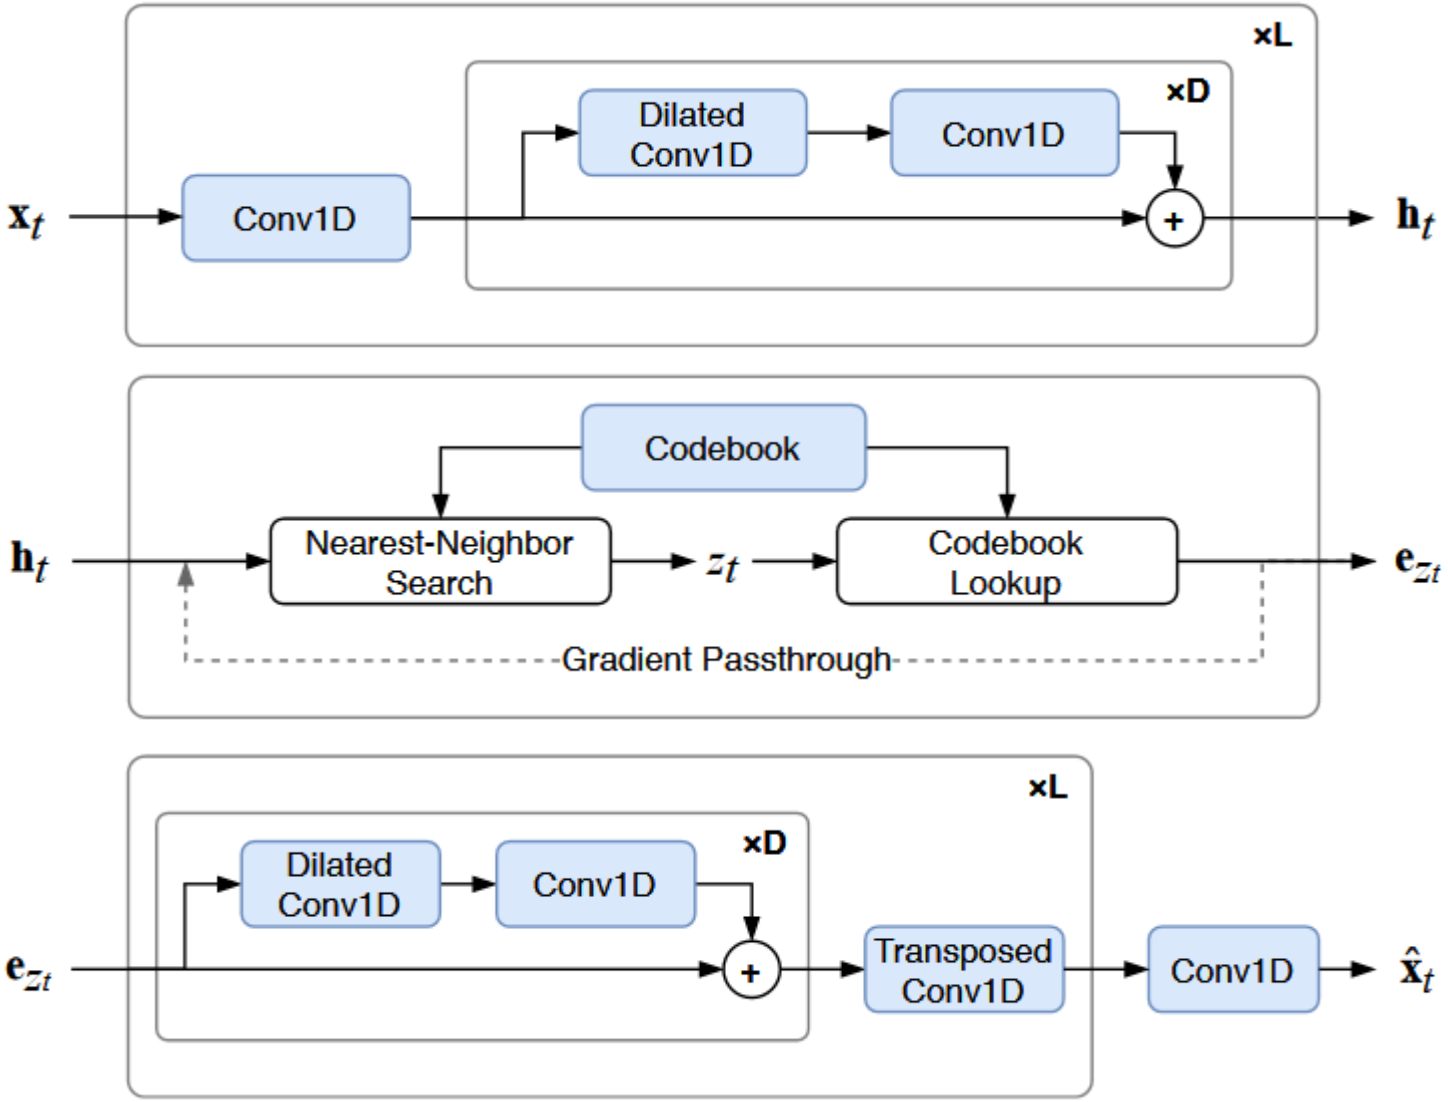

Jukebox: VQ-VAE, Training Hacks

- Random restarts for embeddings + EMA
- Separated Autoencoders
- Spectral Loss (for high-frequencies reconstruction):

$$\mathcal{L}_{\text{spec}} = \| |\text{STFT}(\mathbf{x})| - |\text{STFT}(\hat{\mathbf{x}})| \|_2$$

Jukebox: Learning Prior

$$p(\mathbf{z}) = p(\mathbf{z}^{\text{top}}, \mathbf{z}^{\text{middle}}, \mathbf{z}^{\text{bottom}})$$
$$= p(\mathbf{z}^{\text{top}})p(\mathbf{z}^{\text{middle}}|\mathbf{z}^{\text{top}})p(\mathbf{z}^{\text{bottom}}|\mathbf{z}^{\text{middle}}, \mathbf{z}^{\text{top}})$$

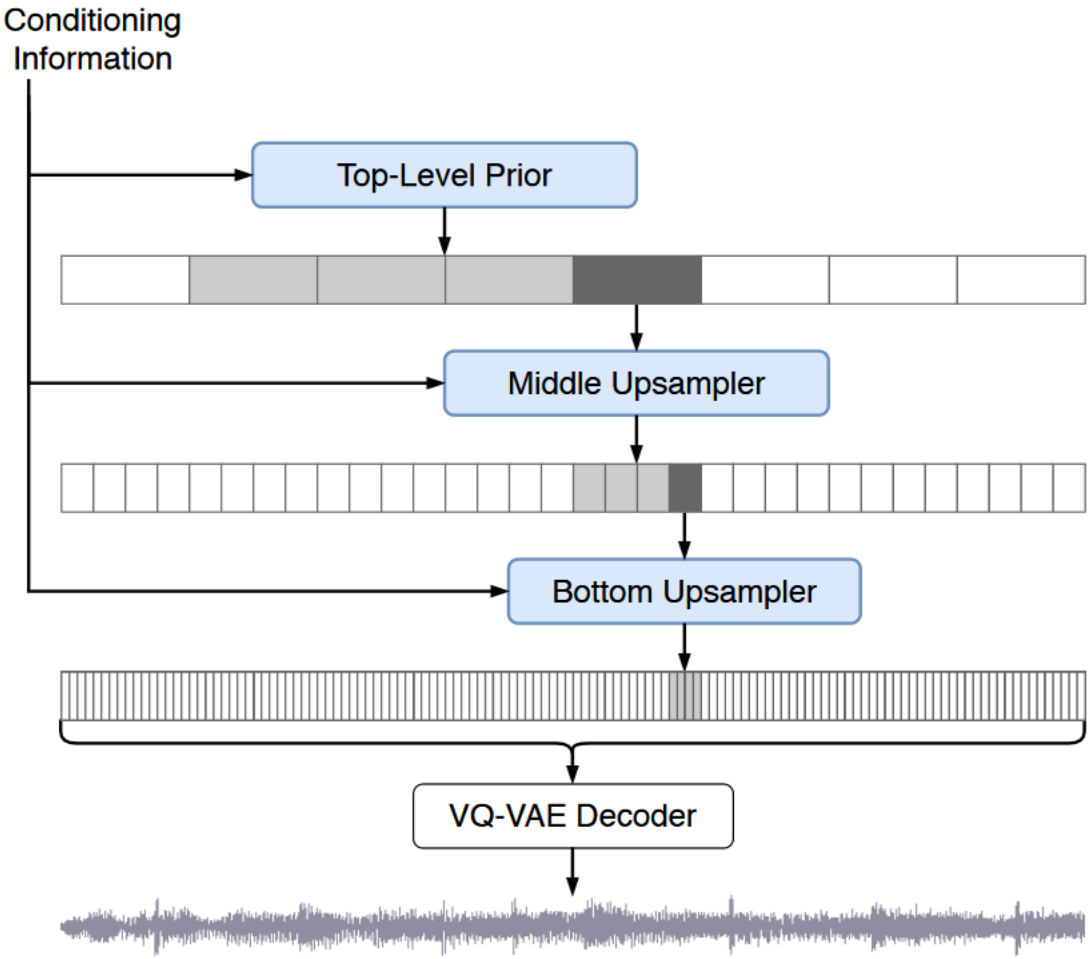

Jukebox: Learning Prior

Top-level prior architecture:

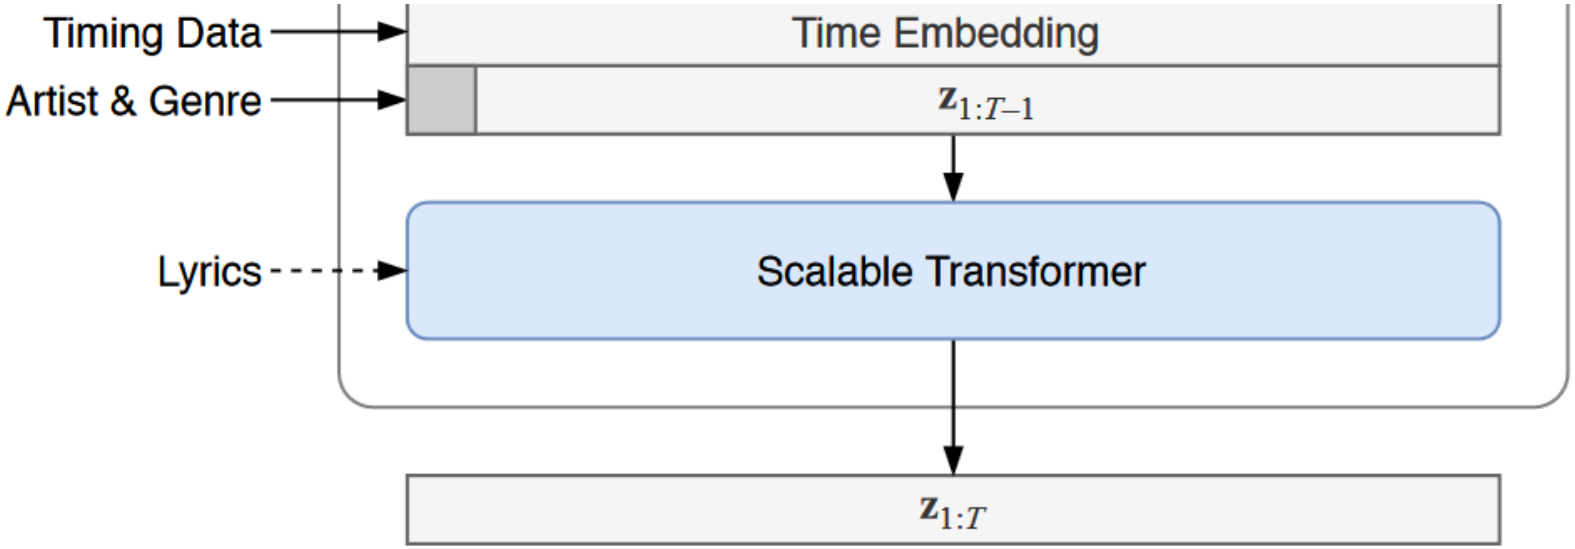

Jukebox: Learning Prior

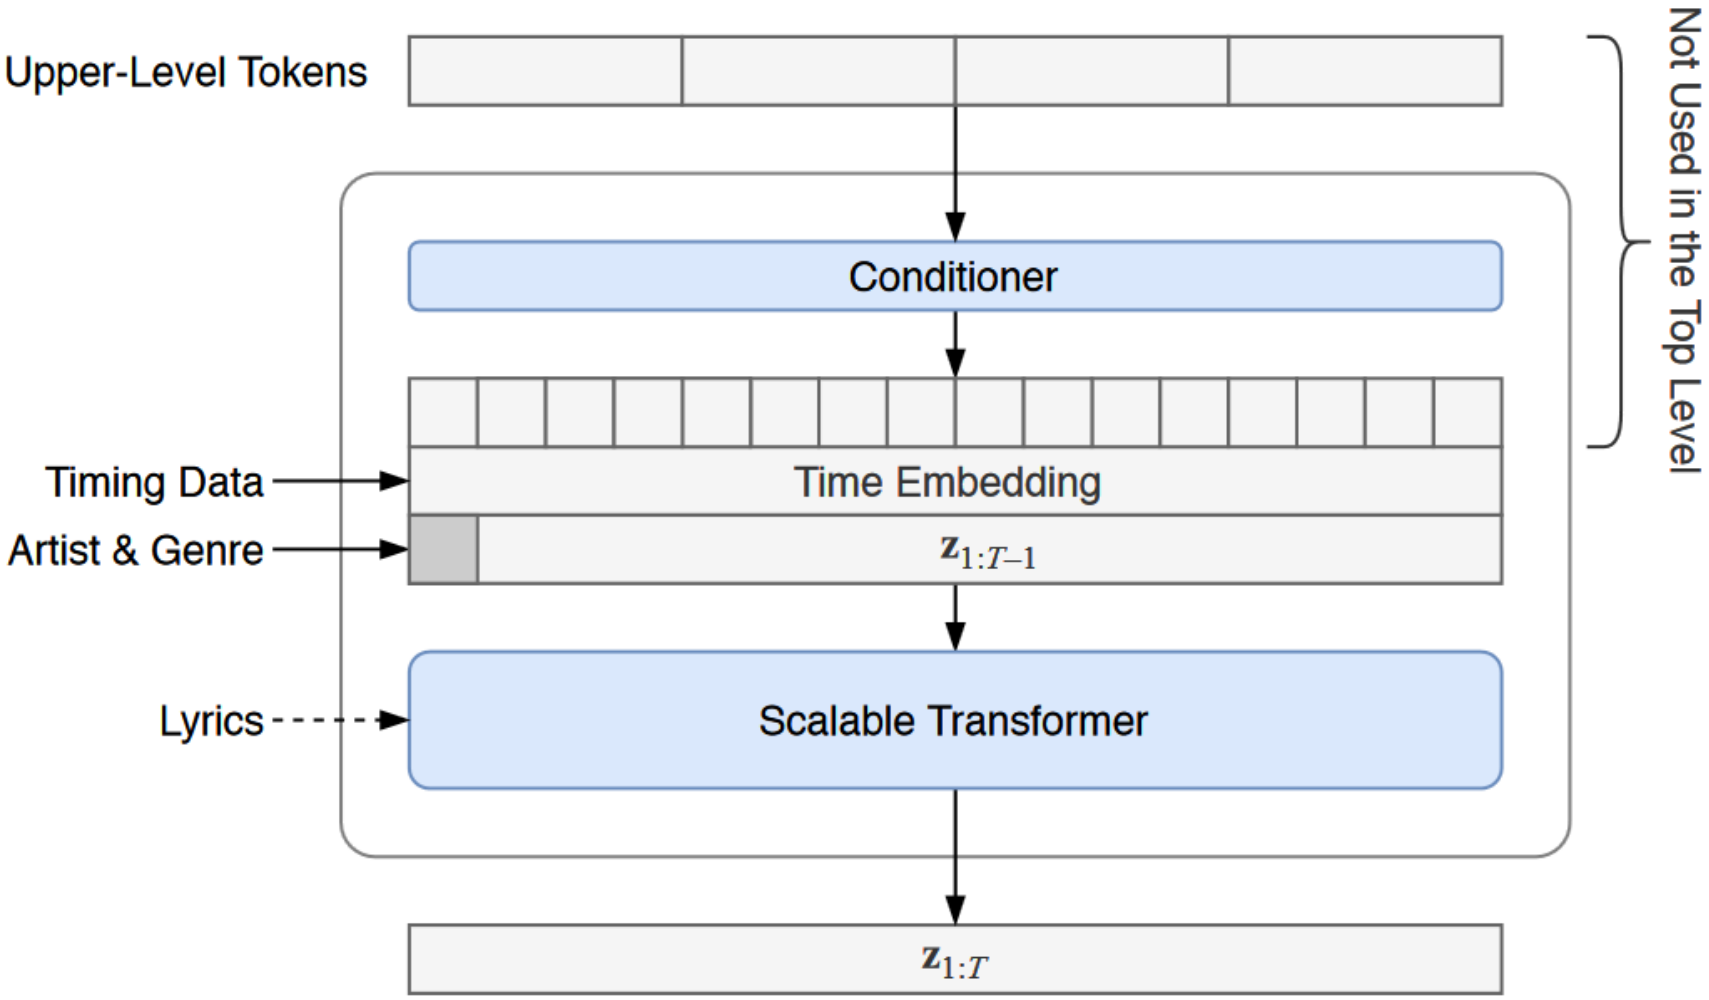

Jukebox: Conditioner

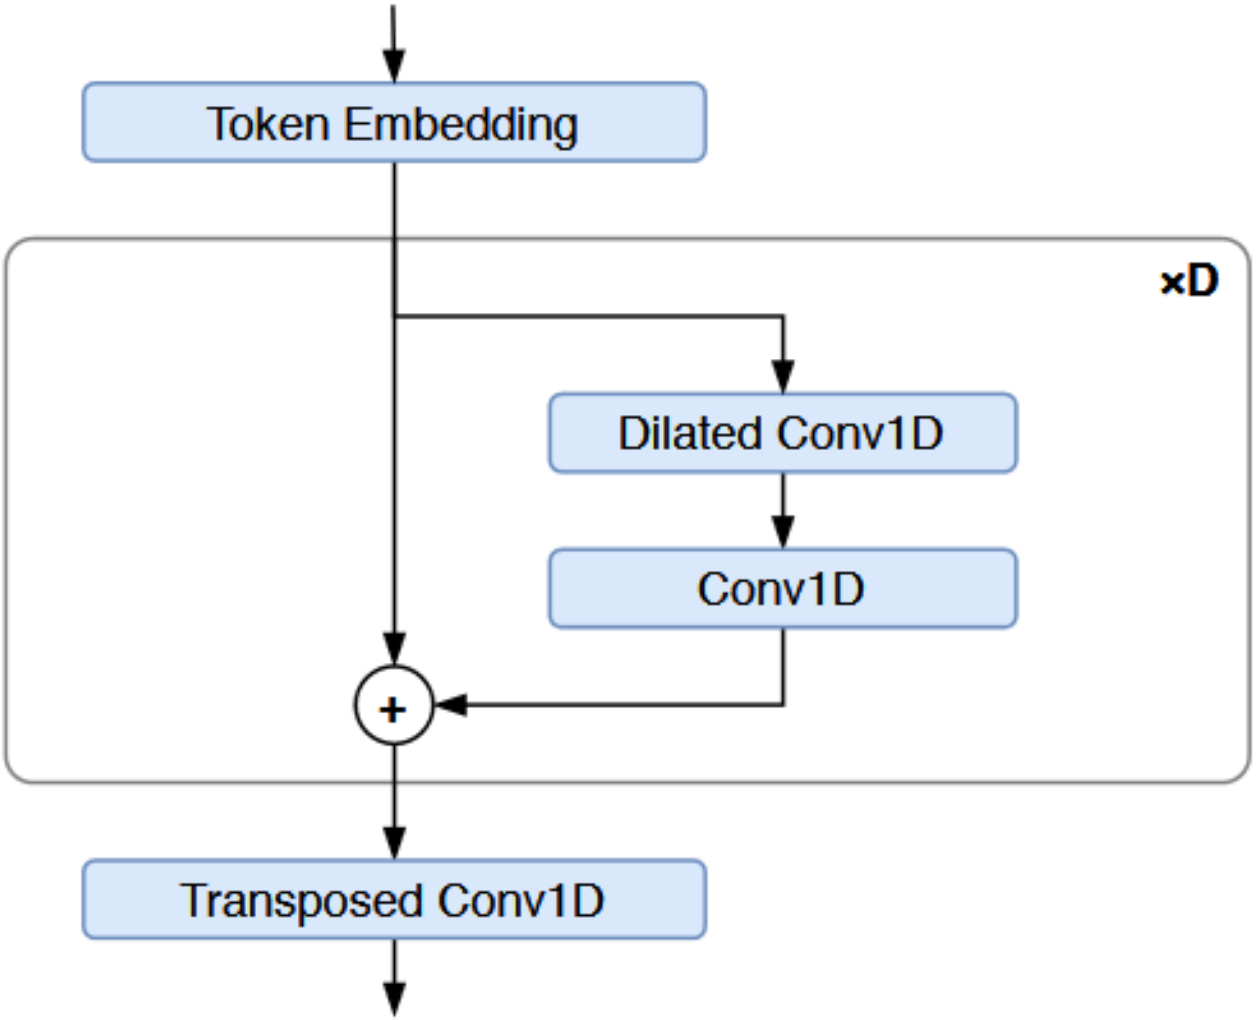

Jukebox: Scalable Transformer

Firstly, let's recall Sparse Transformer:

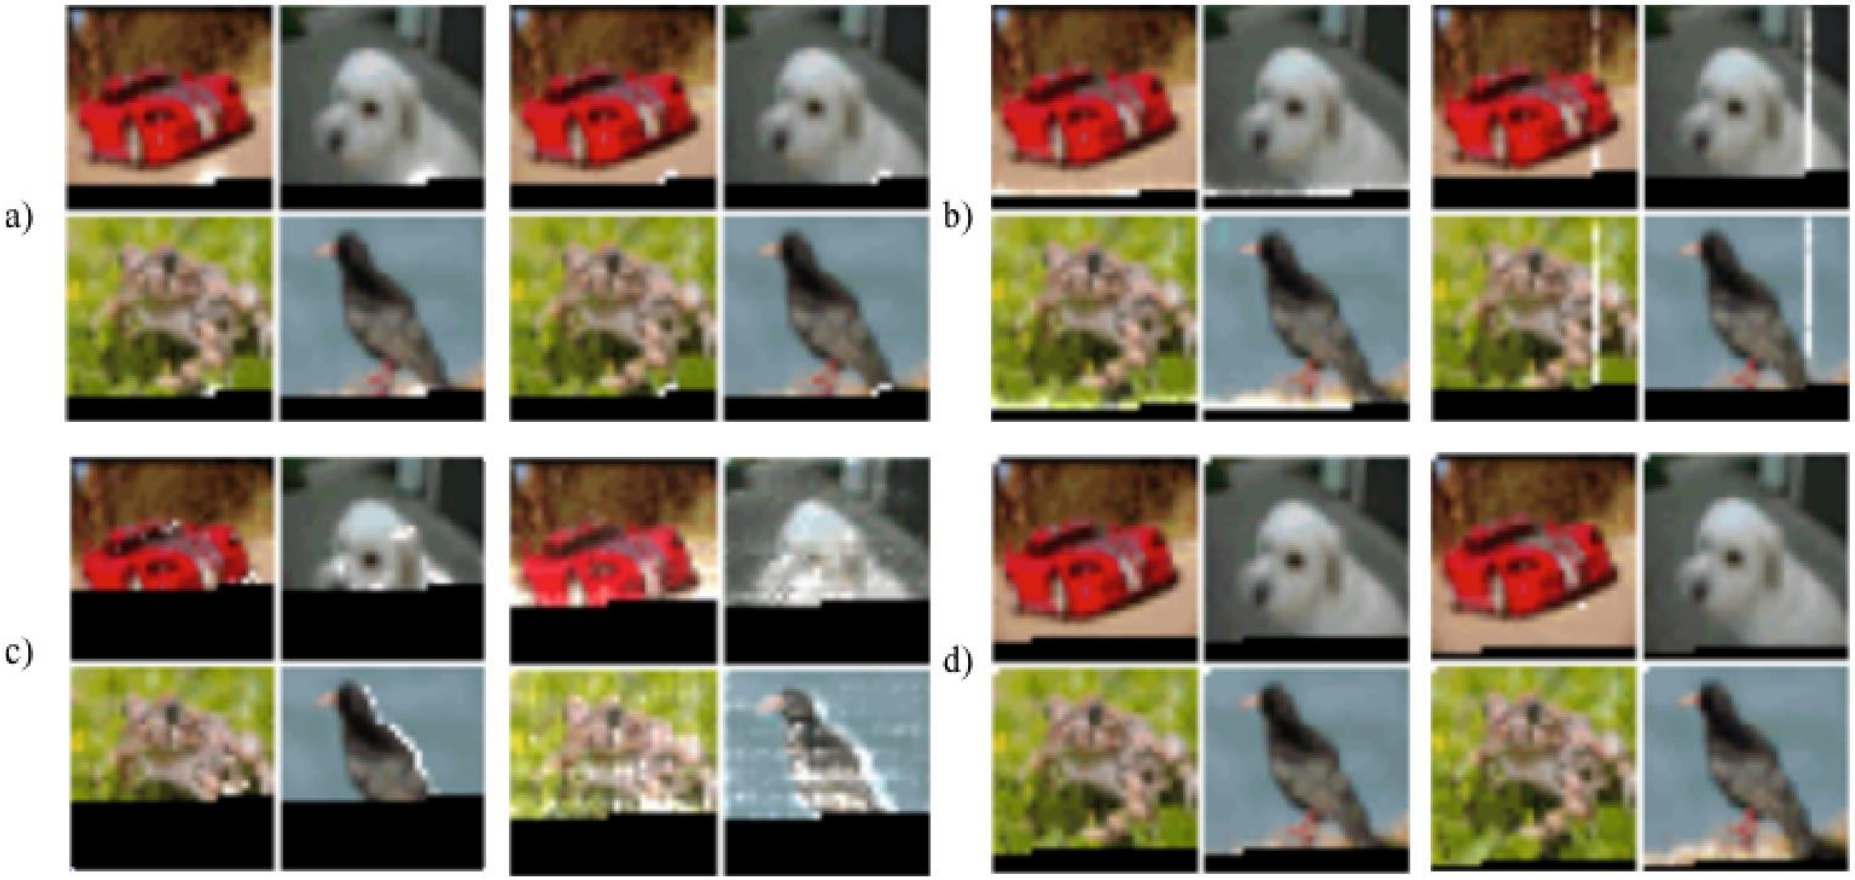

Jukebox: Scalable Transformer

- Vanilla Transformer

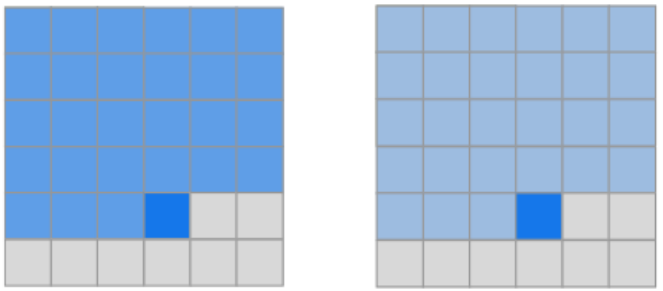

- Sparse Transformer

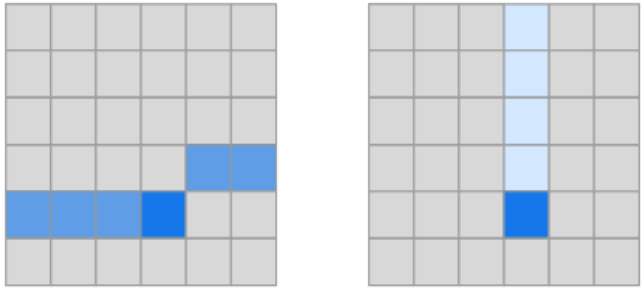

- Scalable Transformer

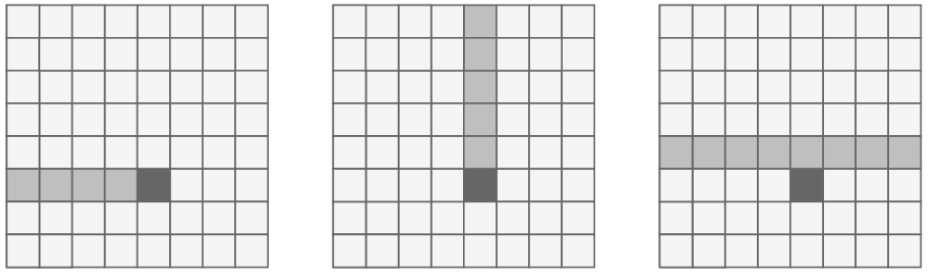

Jukebox: Lyrics Conditioning

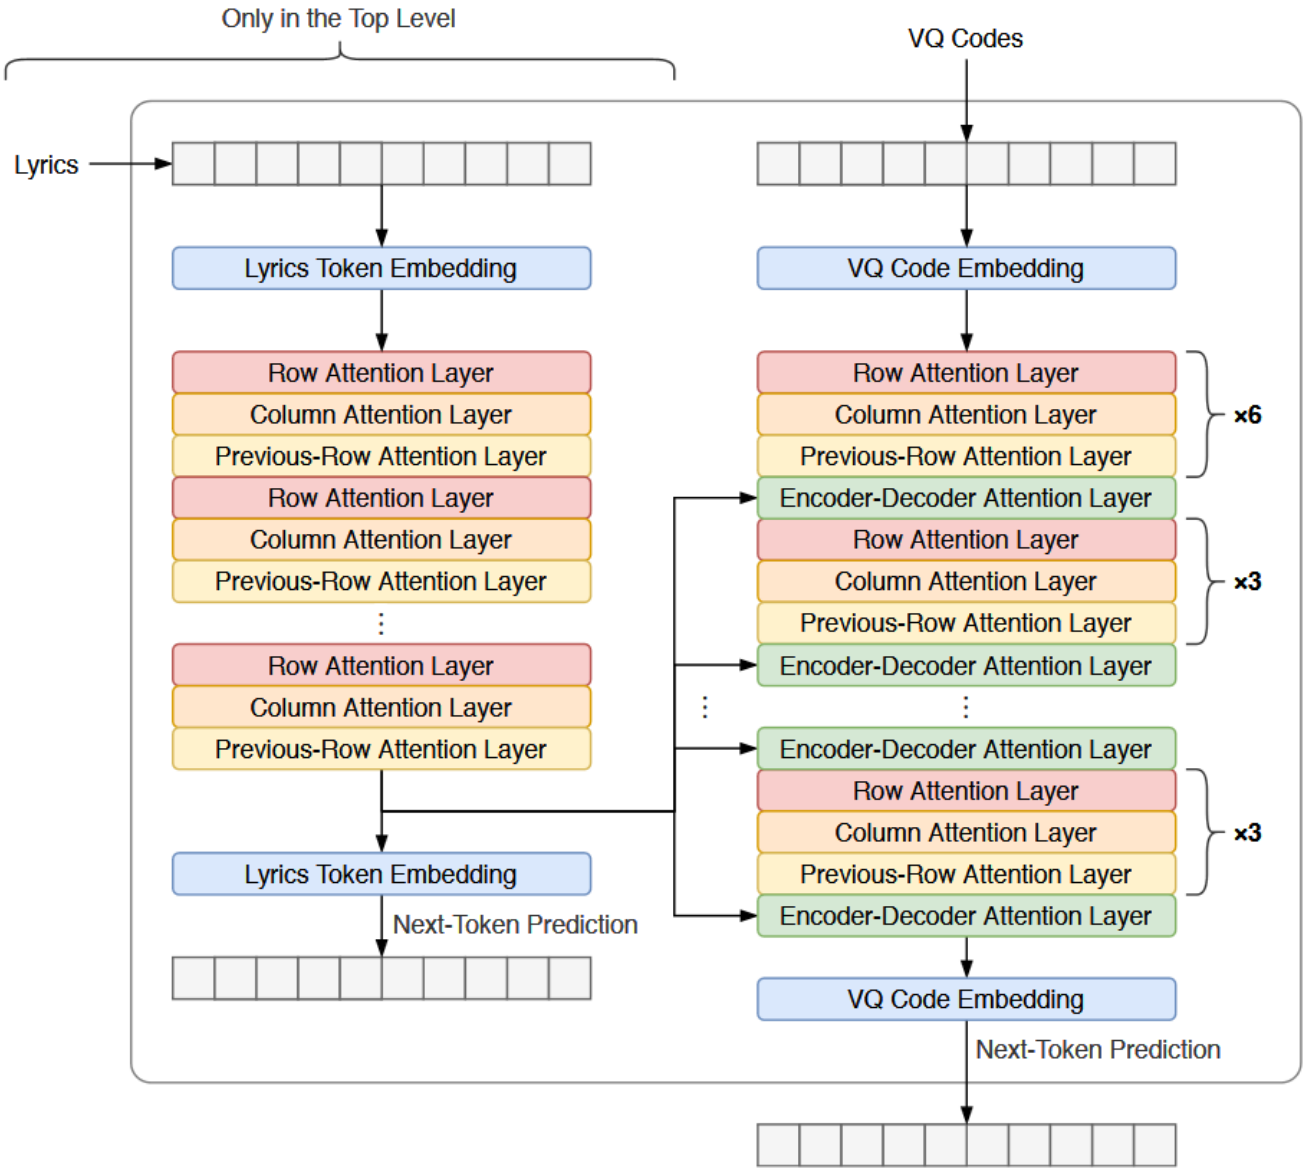

Jukebox: Ablation Study

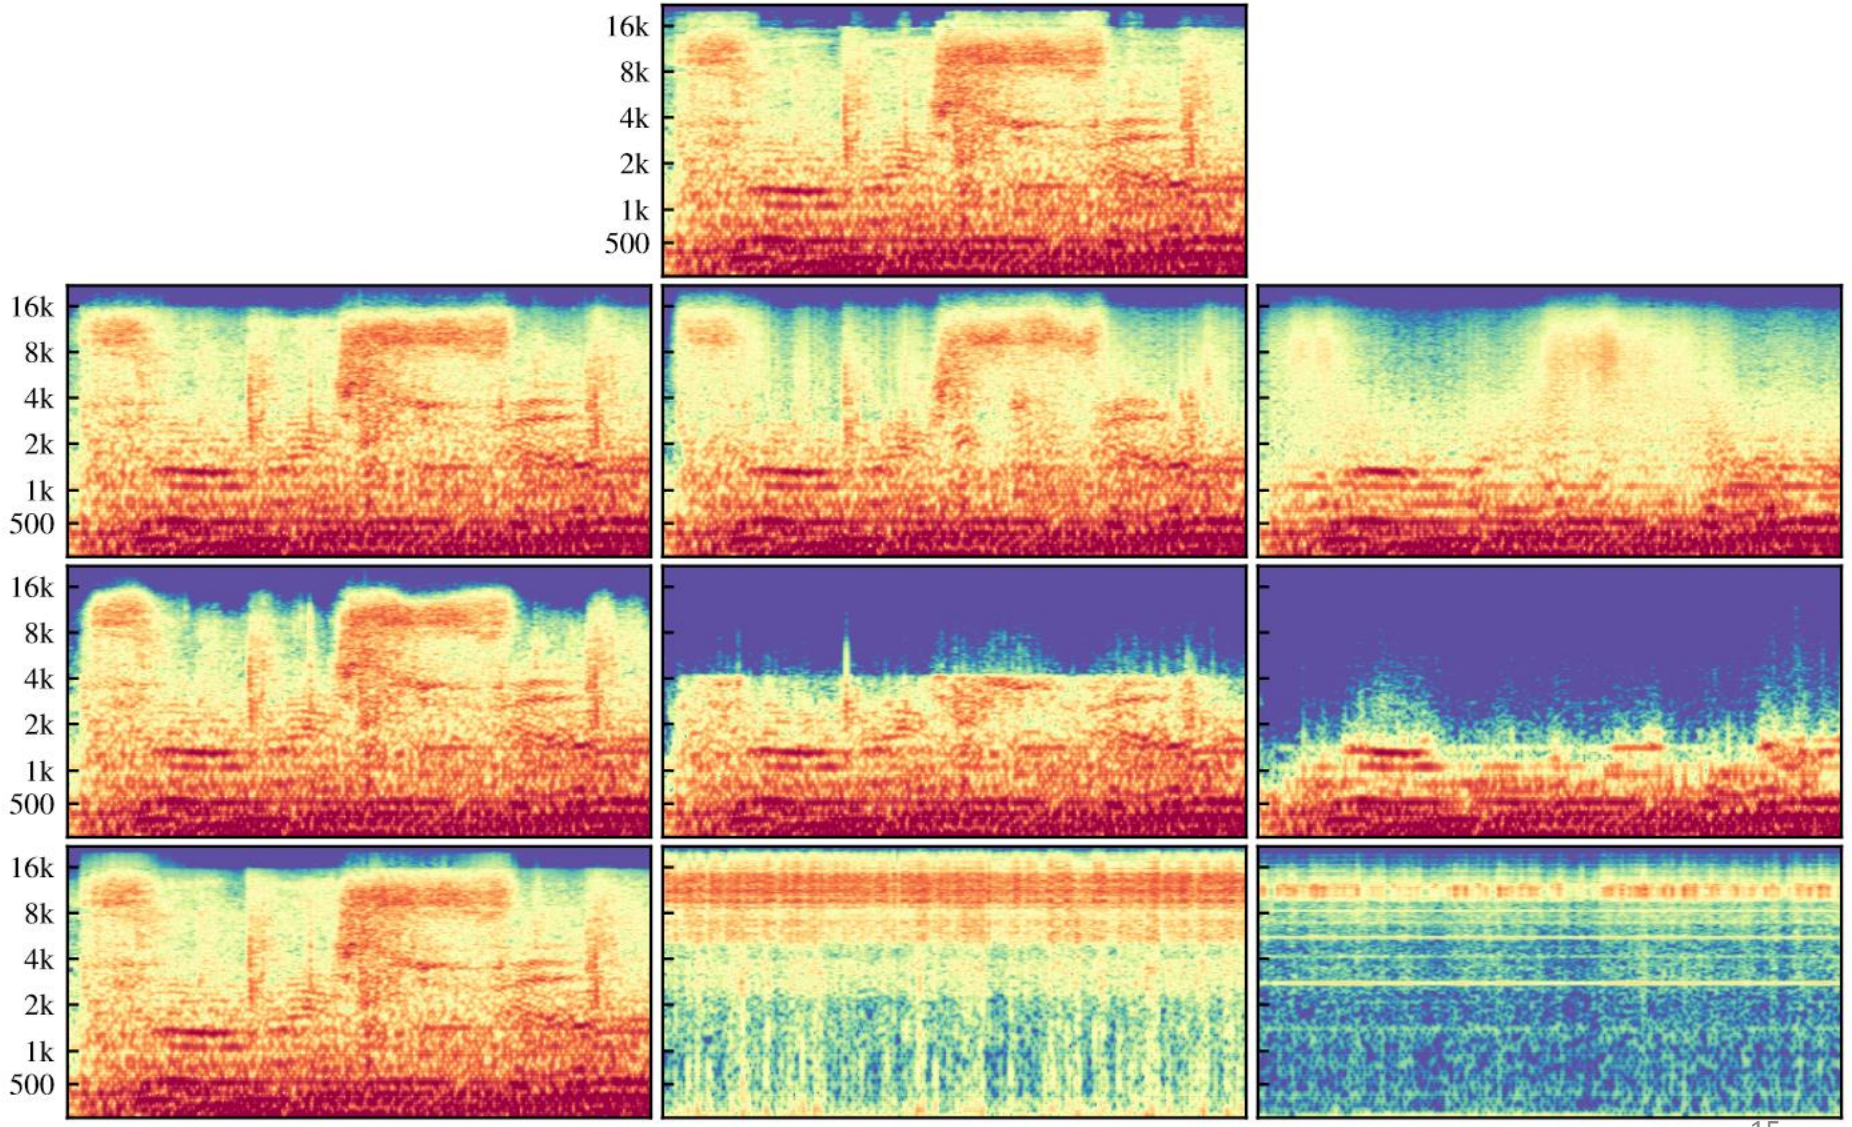

Jukebox: Ablation Study

Level	Hop length	Spectral convergence (dB) Without restart	With restart
Bottom	8	−21.1	−23.0
Middle	32	−12.4	−12.4
Top	128	−8.3	−8.3

Codebook size	Spectral convergence (dB)
256	−15.9
2048	−23.0
No quantization	−40.5

Ablation	Spectral convergence (dB)
None	−8.3
Without spectral loss	−6.3
With single autoencoder	2.9